

Human Natural Structure–Paris AI Action Summit Structural Integrity Report

*A Structural Mapping of Human-Rights-Based AI Governance Requirements to
HNS, SMS-6, EVA, ECS, and Conformance*

Prepared by Satoru Hara
Natural Structure Works (NSW)
July 2026

Status

Independent analytical report. This document does not imply endorsement, adoption, or recognition of HNS by the Paris AI Action Summit, the United Nations, OECD, the Alan Turing Institute, SaferAI, CEN, CENELEC, ITU, ISO, IEC, or any other institution.

Method and evidence note. *This report compares the stated themes of the February 2025 Paris AI Action Summit official side-event summary with the architecture and claims contained in the six-volume Human Natural Structure Series. It distinguishes: (1) source-established governance requirements; (2) HNS engineering proposals; and (3) conclusions that remain hypotheses requiring independent validation.*

Contents

Right-click and select Update Field to refresh the table of contents.

0. Executive Summary

The Paris AI Action Summit side-event summary identifies a recurring implementation problem in global AI governance: human-rights principles are increasingly expected to shape technical standards, yet those principles must still be translated into concrete, testable, auditable, and operational requirements. The summit document repeatedly connects responsible AI to standards, risk management, assurance, certification, investor oversight, corporate due diligence, human supervision, and multistakeholder legitimacy.

Human Natural Structure (HNS) is relevant to this problem because it is designed as an external structural verification architecture rather than as another high-level ethics code. HNS proposes fixed structural coordinates (HNS-36/144/864), contextual grounding through SMS-6, machine-readable audit records through EVA, and intervention or override through ECS. In principle, these components could convert broad governance expectations into observable evidence at the point of AI inference.

The strongest alignment is not that HNS proves human-rights compliance. It does not. Rather, HNS offers a candidate technical substrate for four functions the summit discussion treats as necessary: translating principles into implementation rules, producing traceable evidence, preserving independent oversight, and coordinating obligations across standards and regulatory frameworks.

- High structural alignment: operationalisation, auditability, provenance, human oversight, cross-framework coherence, and vendor-independent verification.
- Partial alignment: substantive human-rights interpretation, impact on affected communities, supply-chain labour rights, investor stewardship, and participatory standard-setting.
- Key limitation: HNS can structure and record evaluations, but it cannot decide by itself which human-rights interpretation is legitimate, who represents affected groups, or whether a standard-setting process is politically inclusive.
- Conclusion: HNS is best positioned as a structural integrity and evidence layer beneath human-rights-based governance—not as a replacement for human-rights law, democratic deliberation, due diligence, or institutional accountability.

1. Purpose, Scope, and Analytical Question

1.1 Purpose

This report evaluates whether the HNS full-stack architecture can respond to the governance and standardisation needs articulated in the official side-event summary, *Shaping the Future of AI Governance: Standards, Risk Management, and Responsible Practices*, issued following the Paris AI Action Summit 2025.

1.2 Primary analytical question

Can HNS supply a defensible structural integrity layer that translates human-rights-informed AI governance requirements into explicit constraints, observable checks, auditable records, and controlled interventions?

1.3 What this report does not claim

- HNS has not been endorsed or adopted by any organisation associated with the summit.
- Technical conformance to HNS would not, by itself, establish legal compliance with the UN Guiding Principles on Business and Human Rights, the OECD Guidelines, the EU AI Act, or any national law.
- This report does not treat HNS’s preliminary PoC results as independent proof of superior accuracy or universal error elimination.
- The report assesses architectural correspondence and implementation potential, not institutional acceptance.

2. Structural Integrity Requirements Emerging from the Summit

Across the two panels, the summit summary can be read as defining a structural integrity agenda for AI governance. “Structural integrity” in this report means that governance commitments remain coherent from principle to technical implementation, from implementation to evidence, and from evidence to accountability.

No.	Requirement	Operational meaning
1	Human rights embedded in technical standards	Fundamental rights must shape technical requirements rather than remain external ethical aspirations.
2	Operational translation	Principles must be documented clearly enough for engineers and programmers to implement.
3	State–business responsibility separation	Public duties and corporate responsibilities must be distinguishable while remaining coordinated.
4	Global coherence	Standards should reduce fragmentation across jurisdictions, industries, and investment portfolios.
5	Risk management across the lifecycle	Impact assessment, due diligence, downstream-use risk, and continuous change must be addressed.
6	Assurance and certification	Standards must support verification, auditing, certification, and legally meaningful evidence.
7	Human oversight and enforceability	Governance must move beyond voluntary ethics toward rules that can trigger review, intervention, and remedy.
8	Multistakeholder legitimacy	Technical experts, human-rights specialists, civil society, government, industry, and affected communities must be able to contribute.
9	Transparency of standard-setting	The process, documents, assumptions, and decision rules must be open enough to challenge.
10	Practical usability	Frameworks must be actionable for companies and startups rather than remaining abstract policy statements.

3. HNS as a Structural Integrity Architecture

3.1 Core architecture

The HNS Series describes a three-axis model. Axis 1 is the statistical or generative model. Axis 2 is an explicit structural and causal constraint plane implemented through HNS-36, HNS-144, and HNS-864. Axis 3 is a grounding and governance plane implemented through SMS-6. The hns36 runtime and EVA are intended to intercept, evaluate, constrain, and log candidate outputs without requiring the verification rules to be absorbed into the base model’s weights.

Component	Integrity function	Governance relevance
HNS-36	Base structural coordinates	Detects unbridged movement between causal layers and abstract categories.
HNS-144	Logical relation extension	Refines concept boundaries and reduces category ambiguity.
HNS-864	Precision causal audit	Tests whether causal statements have a traceable structural path.
SMS-6	Six-layer contextual grounding	Checks somatic, protocol, ecosystem, economic, governance, and universal context.
EVA	External Verification Architecture	Generates an externally inspectable record of checks, corrections, and provenance.
ECS	External Control System	Provides controlled intervention, escalation, blocking, or human override.
LCS	Life Cycle Structure	Extends evidence and controls across design, deployment, operation, incident, and retirement stages.

3.2 The defensible HNS claim

The most defensible current claim is architectural: an external layer using fixed coordinates can produce a machine-readable record of how an AI output was constrained. This auditability is true by design if the implementation faithfully records every check. It is a narrower and more credible claim than asserting that HNS independently guarantees factual truth, legal compliance, or elimination of all AI errors.

4. Structural Mapping: Summit Requirements to HNS

Summit requirement	HNS mechanism	Potential contribution	Boundary / residual gap
Human rights in technical standards	SMS-5; HNS ontology; conformance clauses	Human-rights requirements can be represented as explicit normative constraints and mapped to output evidence.	HNS does not determine the authoritative content of the right.

Summit requirement	HNS mechanism	Potential contribution	Boundary / residual gap
Actionable engineering guidance	HNS-36/144/864; runtime rules	Broad principles can be decomposed into coordinate checks, transition rules, and pass/fail records.	Requires domain-specific profiles and test suites.
Transparency and assurance	EVA; provenance records	Each intervention can be linked to a structural reason, timestamp, model, rule, and output.	Completeness and decision relevance require independent audit.
Human oversight	ECS; escalation pathways	Defined thresholds can trigger pause, rejection, human review, or override.	Human reviewers need authority, competence, and organisational protection.
Risk and impact management	SMS-6; LCS	Signals can be organised across operational, organisational, economic, governance, and systemic layers.	Affected-person impacts need empirical and participatory assessment.
Global coherence	Canonical coordinates; crosswalks	A common evidence model can map multiple standards without replacing them.	Crosswalks may oversimplify legal differences.
Responsible business conduct	EVA + LCS evidence packages	Logs and controls can support due diligence, internal governance, investor engagement, and incident review.	Corporate incentives and remedy remain institutional questions.
Multistakeholder governance	Open profiles, versioned rules, challenge records	Rules and mappings can be exposed for review, amendment, and disagreement logging.	HNS cannot solve financial, political, or expertise barriers to participation.
Startup usability	Layered profiles and minimum controls	A lightweight profile could provide basic structural checks without full enterprise implementation.	Needs usability testing and cost evidence.
Enforceable rules beyond ethics	Conformance criteria; ECS actions	Normative clauses can be connected to test evidence and predetermined intervention.	Legal force comes from law, contract, procurement, or adopted standards—not HNS itself.

5. The HNS Structural Integrity Chain

The principal value of HNS in the Paris governance context is its ability to propose an end-to-end integrity chain. Each link must remain intact; failure at any link breaks the claim that a high-level commitment has been operationalised.

- 1. Normative source** — Human-rights law, regulation, adopted standard, contract, or governance policy.
- 2. Interpretive profile** — A documented interpretation created by competent legal, technical, and affected-community participants.
- 3. Structural encoding** — Mapping of the profile into HNS coordinates, SMS layers, invariants, thresholds, and exceptions.
- 4. Runtime verification** — Independent checks of model inputs, intermediate decisions where available, and outputs.
- 5. Evidence generation** — EVA record containing provenance, rule version, coordinates, reasons, corrections, and unresolved uncertainty.
- 6. Control and escalation** — ECS action: permit, warn, constrain, block, refer to human review, or suspend.
- 7. Organisational accountability** — Named owner, review deadline, incident process, remediation, appeal, and remedy.
- 8. External assurance** — Independent assessment of rule validity, implementation fidelity, evidence completeness, and outcomes.
- 9. Feedback and revision** — Update profiles based on incidents, affected-community input, regulatory change, and audit findings.

6. Analysis of the Summit’s Core Themes

6.1 From human-rights principles to technical requirements

The summit summary’s central implementation demand is translation. HNS is structurally suited to translation because its outputs are based on named coordinates and invariants rather than implicit statistical preference alone. A human-rights requirement can be represented as a profile specifying protected interests, prohibited transitions, evidence thresholds, affected contexts, and escalation conditions. The result could be inspected and challenged by technical and non-technical reviewers.

This suitability is conditional. A coordinate model can preserve the internal consistency of an interpretation but cannot establish that the interpretation is legally correct or socially legitimate. The normative profile must therefore be authored and governed outside the model by qualified institutions and affected stakeholders.

6.2 Audit, assurance, and certification

The summit links standards to assurance and certification. EVA directly addresses this connection by treating evidence as a first-class output. A conforming implementation could produce records for sampling, replay, incident investigation, control testing, and third-party assessment. HNS therefore aligns most strongly with the evidentiary side of governance: showing what rule was applied, when, to which output, and with what result.

6.3 Human oversight and enforceable intervention

Human oversight becomes meaningful only when humans can see why intervention is required and possess the authority to act. HNS and EVA can provide the reason; ECS can define the action. Together they could support graded controls such as warning, constrained regeneration, mandatory approval, automatic blocking, and system suspension. This is more concrete than a general statement that “a human remains in the loop.”

6.4 Responsible business conduct and investor engagement

The summit highlights due diligence, downstream risk, and investor leverage. HNS evidence could support these processes by giving companies and investors a common vocabulary for structural failures, control coverage, unresolved risks, and incident trends. Investors could request evidence-based governance indicators rather than relying solely on policy statements. However, HNS cannot ensure that investors act on the evidence or that companies provide remedy when harm occurs.

6.5 Inclusive global standardisation

HNS may reduce technical participation barriers if its coordinates, profiles, and audit records are published in accessible forms. A shared structural vocabulary can help human-rights experts and engineers discuss the same system. Yet the summit identifies deeper barriers—cost, travel, access to documents, expertise asymmetry, and institutional dominance. These are governance and political-economy problems beyond the reach of an architecture. The appropriate HNS response is procedural support: open documentation, transparent change records, minority opinions, challenge mechanisms, and accessible reference implementations.

7. Proposed Paris Structural Integrity Profile for HNS

A practical next step would be a non-regulatory “Paris Structural Integrity Profile” within the HNS ecosystem. It would not claim official summit status. The name would identify the source themes used to define the profile.

ID	Requirement	Proposed normative statement
PSI-01	Normative traceability	Every technical rule SHALL identify its source principle, legal basis, policy owner, version, and interpretation date.
PSI-02	Rights-impact context	Every high-impact rule set SHALL document affected groups, foreseeable harms, protected interests, and known uncertainties.
PSI-03	Structural encoding	Each requirement SHALL be mapped to HNS coordinates and/or SMS layers with explicit pass, fail, exception, and escalation logic.
PSI-04	Independent evidence	Every verification event SHALL produce an EVA record that is separable from the base model’s self-description.

ID	Requirement	Proposed normative statement
PSI-05	Human intervention	High-severity failures SHALL trigger a defined ECS response and identify an accountable human or institutional owner.
PSI-06	Downstream use	The profile SHALL include foreseeable downstream deployment contexts and misuse pathways.
PSI-07	Multistakeholder review	The profile SHOULD be reviewed by technical, legal, human-rights, domain, and affected-community representatives.
PSI-08	Transparency and challenge	Rules, assumptions, known gaps, dissenting views, and change history SHOULD be accessible to authorised reviewers.
PSI-09	Lifecycle monitoring	The organisation SHALL review incidents, drift, complaints, audit findings, and regulatory changes throughout the lifecycle.
PSI-10	External validation	Claims of effectiveness SHALL be supported by independent, blinded, multi-model, and domain-relevant testing where feasible.

8. Limitations and Threats to Validity

Normative incompleteness: HNS coordinates are not a substitute for human-rights doctrine, jurisprudence, cultural context, or sector-specific law.

Validation limits: The reported HNS PoC is single-model, framework-internal, and not independently graded. It supports feasibility and auditability more strongly than universal effectiveness.

Taxonomy completeness: The five structural failure categories may omit relevant harms such as discrimination, privacy intrusion, labour exploitation, security abuse, or material exclusion unless domain profiles add them.

Implementation fidelity: A vendor could claim HNS conformance while implementing incomplete logging, weak rules, or non-independent verification. Conformance testing must examine actual behaviour and records.

Metric gaming: Once structural scores become targets, organisations may optimise the score without improving real-world outcomes. Outcome measures and affected-person feedback remain essential.

Power and participation: Transparent coordinates do not by themselves redistribute power within standardisation or corporate governance.

Legal uncertainty: Mapping a technical control to a legal article is supportive evidence, not proof of compliance, conformity, or liability protection.

Operational burden: Fine-grained profiles and audit records may impose cost, latency, data-governance, and privacy burdens that require proportionality analysis.

9. Overall Structural Integrity Assessment

Dimension	Assessment	Reason
Principle-to-implementation traceability	Strong potential	Fixed coordinates, profiles, and explicit rules support traceable translation.
Auditability and provenance	Strongest current contribution	EVA is designed specifically to produce external, machine-readable evidence.
Human oversight and intervention	Strong architectural potential	ECS can connect findings to controlled actions and accountable review.
Cross-standard coherence	Strong potential, subject to mapping quality	HNS can serve as a common evidence layer without replacing source instruments.
Substantive human-rights compliance	Partial	Requires authoritative legal interpretation, impact assessment, and institutional accountability.
Inclusive standard-setting	Limited to supportive tooling	Open rules and records help, but participation barriers are external.
Empirical proof of effectiveness	Not yet sufficient	Independent multi-model and domain validation is still required.
Readiness for adoption as an official standard	Premature	The architecture is a proposed specification and requires external review, pilots, and consensus.

10. Recommendations

1. Publish a short Paris Structural Integrity Profile as a clearly labelled HNS proposal, not as an official summit document.
2. Create a machine-readable crosswalk linking each profile requirement to UNGPs, OECD responsible-business-conduct expectations, relevant EU AI Act obligations, and recognised standards—while distinguishing legal obligation from technical support.
3. Develop at least two domain pilots, preferably one high-impact public-service use case and one enterprise deployment, with independent evaluators and affected-community participation.
4. Make EVA records interoperable through open provenance formats and publish a minimal audit-record schema.
5. Define independence criteria for the external verifier, including separation of control, rule ownership, logging, and audit access from the model provider.
6. Add explicit rights-impact modules covering discrimination, privacy, accessibility, labour, security, and remedy rather than relying only on the five structural hallucination categories.

7. Commission blind, multi-model tests comparing HNS against active controls such as structured reasoning prompts, conventional guardrails, and independent rule engines.
8. Design a lightweight startup profile so that practical implementation does not require the full enterprise stack.
9. Introduce public change logs, dissent records, and challenge procedures to demonstrate that the profile itself is governable.
10. Position HNS consistently as evidence infrastructure for governance, not as an autonomous source of legal or ethical authority.

11. Conclusion

The Paris AI Action Summit side-event summary describes a field moving from broad ethical commitments toward standards, risk controls, assurance, and enforceable accountability. Its unresolved question is how human-rights principles become technical practice without losing their normative meaning or democratic legitimacy.

HNS offers a serious architectural response to one part of that question. Its explicit coordinate system, contextual grounding, external verification records, and control pathways could create a continuous chain from requirement to evidence and intervention. This is a meaningful structural contribution because many current governance frameworks specify what organisations should achieve while leaving the inference-level evidence mechanism underdeveloped.

The correct conclusion is therefore bounded but significant: HNS is structurally compatible with the principal operational needs articulated at the Paris event, especially traceability, auditability, assurance, human oversight, and cross-framework coherence. It is not a substitute for human-rights law, participatory governance, institutional legitimacy, or independent empirical validation. Developed with those boundaries intact, HNS could serve as a candidate structural integrity layer for responsible AI governance.

References

- UN B-Tech Project, AI Standards Hub, SaferAI, and OECD. (2025). Shaping the Future of AI Governance: Standards, Risk Management, and Responsible Practices — Key Takeaways from the official side-event to the Paris AI Action Summit 2025.
- Hara, S. (2026). Human Natural Structure Origin & Theoretical Foundation. Volume 1, Human Natural Structure Series. Natural Structure Works.
- Hara, S. (2026). Human Natural Structure Foundational & Observational Specification. Volume 2, Human Natural Structure Series. Natural Structure Works.
- Hara, S. (2026). Human Natural Structure Analytical Specification. Volume 3, Human Natural Structure Series. Natural Structure Works.
- Hara, S. (2026). Human Natural Structure Social Meta-Structure Specification. Volume 4, Human Natural Structure Series. Natural Structure Works.
- Hara, S. (2026). Human Natural Structure External Verification Architecture. Volume 5, Human Natural Structure Series. Natural Structure Works.

Hara, S. (2026). Human Natural Structure Standards & Conformance Specification. Volume 6, Human Natural Structure Series. Natural Structure Works.

United Nations. Guiding Principles on Business and Human Rights.

OECD. Guidelines for Multinational Enterprises on Responsible Business Conduct.