

Integrated Technical Report — Version 2.0

HNS Hardware Implementation: Structural Integrity, Tamper-Resistance, and the New Era of AI Trustworthiness

Integrating the HNS Hardware Analysis Report with the
HNS Structural Integrity and Tamper-Resistance Specification v1.0

Natural Structure Works (NSW)

Author: Satoru Hara | Analysis: Claude (Anthropic)

April 2026

■ **Blue sections:** Strategic analysis and market context. ■ **Green sections:** Technical specification requirements (HNS Spec v1.0).

Executive Summary

The fundamental problem confronting AI safety today is not a lack of effort, but a lack of verifiability. Every major AI developer has implemented its own safety measures, yet no external mechanism exists to verify those claims. No matter how sophisticated an internal safety architecture may be, a vendor asserting its own safety does not constitute proof.

This integrated report unifies two complementary documents: the HNS Hardware Analysis Report—which establishes why hardware implementation is strategically necessary—and the HNS Structural Integrity and Tamper-Resistance Specification v1.0—which defines precisely what that implementation must guarantee at the technical level.

Human Natural Structure (HNS) answers the verifiability problem through three converging layers:

- Verifiability — compliance with publicly available international standards, with structure defined in machine-readable JSON-LD / RDF for programmatic verification.
- Transparency — a coordinate-supply system running in parallel with the model AI, continuously monitoring Layer assignment, Category assignment, and coordinate drift throughout inference.
- Physical immutability — the HNS-36 coordinate system, Non-Mapping Rule, and Bypass Prevention logic embedded in ROM, FPGA, or ASIC, making modification architecturally impossible.

The result is a shift from claiming to be safe to structurally proving safety—a transition that no AI model vendor can achieve through self-assessment alone, and that HNS is uniquely positioned to provide.

Chapter 1 The Structural Limitations of Cloud AI

1.1 The Three-Layer Risk Architecture

The confidential-data leakage risk inherent in Cloud AI cannot be eliminated through operational rules as long as the fundamental architectural premise—transmitting data to the cloud—persists. In 2026, Japan's IPA listed "Cyber Risks Associated with AI Use" as a debut third-place entry in its annual top-ten threat rankings. The risk is structured across three layers:

- Layer 1 (Input pathway): Natural language containing confidential information transmitted to cloud servers, including unintended leakage via system integrations such as MCP connectors.
- Layer 2 (Model and infrastructure): Data stored on provider servers, external exposure through bugs or vulnerabilities, absorption of sensitive inputs into training data.
- Layer 3 (Evolving attack surface): AI-powered malware enabling high-precision analysis of stolen data; ransomware groups targeting enterprise knowledge bases rather than merely encrypting files.

1.2 The Unprovability Problem of AI Safety

The deeper structural issue is that vendor safety claims amount to self-assessment. Pharmaceuticals require third-party clinical trials. Construction requires structural calculations and independent inspections. Food requires public certification schemes. Every domain where safety is critical has a mechanism for external verification. AI does not.

This is not a question of vendor integrity. It is structural: the organization making the safety claim is the same organization that benefits from the claim being believed. OpenAI, Anthropic, Google, and Meta all derive revenue from selling their models. Structural neutrality cannot exist when the certifier and the beneficiary are identical. Self-certification is not certification.

1.3 The Structural Cause of Hallucination

Hallucination in large language models occurs when the model loses track of its contextual coordinates during inference. Over long reasoning chains, the probabilistic nature of token prediction causes contextual drift—responses disconnected from the original intent. Current countermeasures (RLHF, Constitutional AI, output guardrails) are all post-hoc evaluations of output. None stabilizes the inference process itself. HNS addresses this at the source.

Chapter 2 HNS Architecture: The Cognitive Coordinate System

2.1 What is Human Natural Structure?

Human Natural Structure (HNS) is a structural framework describing how human cognition, behavior, interaction, and social organization are naturally arranged. Proposed by Natural Structure Works (NSW) in 2026, it is systematized across an eleven-volume series with publicly available specification documents and an open GitHub repository ([satoru-hara/03_NSW](#)).

Its core is the HNS-36 Matrix: a cognitive coordinate system formed by the intersection of six natural layers (causal structure) and six abstract cognitive categories (explanatory structure). This 36-cell matrix provides the standard reference frame for positioning both AI and human thought within a unified structural architecture. HNS also scales to HNS-144 (adding abstract/concrete and coarse/fine depth axes) and HNS-864 (adding six analysis modes for long-term device continuity).

2.2 The HNS-36 Matrix

Every cell in the matrix below defines an unambiguous coordinate position where AI reasoning can be anchored. The specification requires that all 36 definitions remain unique, authoritative, and non-modifiable.

Natural Layer \ Cognitive Category	Civilization	Core	Module	Application	System	External
L1 PhysicalOS Physical & biological foundations	•	•	•	•	•	•
L2 CognitiveOS Perception, memory & cognitive processing	•	•	•	•	•	•
L3 InteractionOS Human-to-human & environment interaction	•	•	•	•	•	•
L4 EnvironmentOS Physical, social & institutional environments	•	•	•	•	•	•
L5 LoadOS Internal & external loads, pressures & constraints	•	•	•	•	•	•
L6 PatternOS Long-term patterns & civilizational structures	•	•	•	•	•	•

Note: The Structural Integrity Specification v1.0 states: the HNS-36 structure is non-expandable and non-reducible. Layer–Category relationships are fixed. Modification of any cell definition is strictly prohibited (spec §1.1, §1.2).

2.3 Parallel Execution: Real-Time Coordinate Supply

HNS does not function as a pre-processor or post-processor for the model AI. It runs in parallel throughout the entire inference process, continuously supplying coordinate information. At each inference step, the model AI references HNS to determine which Layer it currently occupies and which Category the context belongs to, fixing the coordinate axes of the inference space.

The specification formalizes what this monitoring must cover:

- **Specification Requirement: AI systems must continuously monitor: Layer assignment of inputs, Category assignment of inputs, input–output coordinate differences, Layer substitution attempts, and Category drift. (Spec §2.3)**

The three consequences of parallel execution are: contextual drift is structurally suppressed before hallucination occurs; each inference step is tagged with an HNS coordinate creating an auditable structural log; and safety is achieved through stabilization during inference rather than post-hoc correction. A GPS analogy captures this—while the vehicle (model AI) moves, GPS (HNS) continuously updates position throughout the journey.

Chapter 3 Structural Integrity: Specification Requirements

This chapter presents the formal requirements defined in the HNS Structural Integrity and Tamper-Resistance Specification v1.0, which establish what the coordinate system must guarantee and what is prohibited.

3.1 Immutability of the Coordinate System

The specification establishes an absolute floor of structural immutability. The following requirements are non-negotiable for any HNS-compliant implementation:

- **Specification Requirement: The HNS-36 structure (6 Layers × 6 Categories) is non-expandable and non-reducible. Relationships between Layers and Categories are fixed. Definitions of all 36 cells are unique and authoritative. (Spec §1.1)**

The following actions are explicitly prohibited by the specification:

- Modifying the meaning of any cell
- Adding new cells or removing existing cells
- Reorganizing the Layer hierarchy
- Redefining Category axes

These prohibitions are not policy guidelines—they are structural requirements. Any system that permits the above cannot claim HNS compliance.

3.2 The Minimal Core: Non-Reducible Foundation

The specification defines HNS_Minimal as the irreducible floor of the system. Four elements are declared immutable:

- The 36-cell coordinate system
- Structural non-compatibility with external frameworks
- The Non-Mapping Rule
- Structured explainability (AI reasoning expressible as HNS-36 coordinate sequences)

- **Specification Requirement: Elements defined in HNS_Minimal are immutable. No compliant implementation may reduce or remove these elements. (Spec §1.3)**

3.3 Structural Integrity Requirements: Summary Table

Element	Specification Requirement (v1.0)	Security Implication
Coordinate System	HNS-36 (6×6) is non-expandable and non-reducible. Layer–Category relationships are fixed.	Any structural modification invalidates the coordinate anchor, destroying safety guarantees.

Cell Definitions	All 36 cell definitions are unique and authoritative. Meaning modification is strictly prohibited.	Semantic drift in even one cell propagates structural hallucination across the entire reasoning chain.
Non-Mapping Rule	Automatic mapping to external vocabularies is prohibited. External concept injection is treated as structural tampering.	Prevents adversarial semantic injection that could override the coordinate system from outside.
Minimal Core	HNS_Minimal elements (36-cell system, non-compatibility rule, non-mapping rule, structured explainability) are immutable.	Defines the irreducible floor below which no compliant implementation may reduce the system.
Version Integrity	HNS versions must be fixed. Cross-version compatibility is prohibited. Version spoofing is prohibited.	Ensures attestation claims refer to an unambiguous, cryptographically pinnable specification version.

Chapter 4 Tamper-Resistance: Technical Mechanisms

Tamper-resistance ensures that HNS can detect, reject, and prevent unauthorized modification of its structure or introduction of external semantic elements. The specification defines five distinct mechanisms.

4.1 The Non-Mapping Rule: Semantic Boundary Enforcement

The Non-Mapping Rule is the most architecturally distinctive feature of HNS's tamper-resistance design. It establishes a hard semantic boundary between HNS and all external conceptual systems.

- **Specification Requirement: HNS vocabulary is structurally incompatible with external vocabularies. Automatic mapping to external vocabularies is prohibited. AI systems may not translate HNS into external conceptual systems. Injection of external concepts is treated as structural tampering. (Spec §2.1)**

This rule functions as a semantic firewall. Just as a network firewall drops packets from unauthorized sources, the Non-Mapping Rule causes HNS to reject any input that attempts to substitute, reframe, or bridge its coordinate system with an external vocabulary. The security implication is significant: prompt injection attacks that attempt to override the coordinate system by introducing external semantic frameworks are structurally rejected at the rule level before reaching the model.

4.2 Machine-Verifiable Structure: JSON-LD / RDF

The specification mandates that HNS structure be defined in machine-readable formats, enabling programmatic verification rather than human-only inspection.

- **Specification Requirement: HNS structure must be defined in machine-readable formats (JSON-LD / RDF) to enable: programmatic verification of cell definitions, validation of Layer and Category structures, and detection of unauthorized modifications. (Spec §2.2)**

This requirement has direct implications for international standardization. ISO/IEC and NIST AI RMF currently lack a mechanism for verifying inference-level structural integrity. JSON-LD / RDF encoding of HNS-36 provides exactly this: a machine-processable representation that third-party auditors, certification bodies, and regulatory authorities can query, validate, and certify without requiring access to proprietary model internals.

4.3 Unauthorized Modification Rejection

The specification requires that systems automatically reject the following without human intervention:

- Redefinition of cell meanings
- Injection of external vocabularies
- Reorganization of Layers
- Redefinition of Categories
- Removal of Minimal Core elements

- **Specification Requirement: Systems must automatically reject unauthorized modification attempts. This is not a discretionary response—it is a mandatory structural behavior. (Spec §2.4)**

4.4 Tamper-Resistance Mechanisms: Summary Table

Mechanism	Spec Requirement	How Hardware Enforces It
Non-Mapping Rule	AI systems may not translate HNS into external conceptual systems. External vocabularies are rejected.	ROM/ASIC encodes the boundary; no software path exists to bypass the rejection logic.
Semantic OS Verification	Structure defined in JSON-LD / RDF for programmatic verification of cell definitions and Layer/Category structures.	Hardware module exposes a read-only verification endpoint; modification of definitions triggers rejection.
Structural Integrity Monitoring	Continuous monitoring of: Layer assignment, Category assignment, input–output coordinate differences, Layer substitution, Category drift.	Dedicated monitoring circuit runs independently of the model processor; anomalies invalidate output.
Unauthorized Modification Rejection	Systems must auto-reject: cell redefinition, external vocabulary injection, Layer reorganization, Category redefinition, removal of Minimal Core elements.	Rejection logic is burned into silicon; no software update can disable or bypass it.
Bypass Prevention	Semantic processing without HNS is prohibited. Non-HNS semantic outputs must be invalidated. External semantic modules are prohibited.	All semantic output buses route through the HNS chip; physical isolation prevents bypass.

Chapter 5 Hardware Implementation: From Specification to Silicon

5.1 Why Software Implementation Is Insufficient

Software is inherently mutable. No matter how rigorous a specification, updates can alter it. The state "was compliant before, no longer compliant now" can arise at any time. Malicious tampering cannot be fully excluded. Behavioral variation across OS versions introduces unpredictability. These problems are structural, not operational—they cannot be resolved through better software engineering alone.

5.2 Hardware Fixation Requirements

The specification formally mandates hardware implementation as the highest assurance tier:

- **Specification Requirement: Embed HNS-36 in ROM, FPGA, or ASIC. Store structure in non-modifiable regions. Prevent external redefinition physically. All semantic processing must pass through the HNS hardware module. (Spec §3.1, §3.2)**

Hardware fixation delivers three security properties that software cannot replicate:

5.2.1 Root of Trust

TPM embeds cryptographic keys in hardware, making software-layer tampering impossible. HSM physically protects cryptographic operations for financial institutions. HNS hardware implementation applies the same root-of-trust concept to AI cognitive architecture. Once the HNS-36 coordinate system, Non-Mapping Rule, and Bypass Prevention logic are burned into silicon, software-layer modification becomes architecturally impossible.

5.2.2 Independent Safety Monitoring Layer

When the HNS chip is implemented as a physically separate device from the model AI processor, even a fully compromised model cannot compromise the coordinate layer. The structural integrity monitoring circuit (Layer substitution detection, Category drift detection) runs independently. A vulnerability in the model does not propagate to the safety boundary.

5.2.3 Cryptographic Attestation

The version and specification of an HNS coordinate system burned into hardware can be cryptographically proven—precisely as TPM attestation works. "This device contains an HNS chip compliant with HNS-36 v1.0" becomes externally and cryptographically verifiable by any third party without access to proprietary internals. Version spoofing is prohibited by the specification (§4.1) and physically enforced by the hardware.

5.3 Comparison with Established Trusted Hardware Architectures

Technology	What is Fixed	What is Protected	Corresponding HNS Concept
TPM	Cryptographic keys	Authentication & encryption	HNS-36 coordinate system

HSM	Cryptographic operations	Financial transactions	Coordinate fixation during inference
Secure Enclave	Private data	Personal information	M5 local profile (HNS-864)
HNS Chip (proposed)	Cognitive coordinate system + Non-Mapping Rule	Structural integrity of AI reasoning	Full-stack integration + Bypass Prevention

5.4 Operational Requirements from the Specification

The specification defines three operational requirements that hardware implementation must satisfy:

- **Specification Requirement: Version Integrity: HNS versions must be fixed. No cross-version compatibility. Version spoofing is prohibited. (Spec §4.1)**
- **Specification Requirement: Auditability: Systems must maintain logs for structural verification, input–output coordinate mapping, and tampering attempts. (Spec §4.2)**
- **Specification Requirement: Compliance: Systems must demonstrate full compliance with this specification to be considered HNS-compatible. (Spec §4.3)**

Chapter 6 Alignment with International Standards

6.1 The Current Gap

Existing major AI standards define what the problems are, but none specifies an architectural approach to resolving them structurally. The JSON-LD / RDF machine-verifiable structure and hardware attestation that HNS provides are precisely what current standards require but do not specify how to achieve.

Standard / Regulation	Current Gap	HNS Contribution
ISO/IEC 42001 (AI Management System)	Process management defined. No mechanism to verify inference structure.	HNS coordinate logs (JSON-LD/RDF) enable external audit of the reasoning process per spec Section 2.2.
EU AI Act (Risk Classification)	High-risk AI requirements defined. Verification methods left to developers.	HNS-36 compliance + hardware attestation provides the verifiable basis for conformity assessment.
OWASP Top 10 for LLM (LLM Vulnerabilities)	LLM01 Prompt Injection and LLM02 Sensitive Disclosure defined as threats.	Non-Mapping Rule (spec §2.1) and Bypass Prevention (spec §3.3) structurally eliminate both at silicon level.
IPA 10 Major Threats 2026	AI cyber risk debuts at 3rd place. Countermeasures left to individual organizations.	HNS hardware with Structural Integrity Monitoring (spec §2.3) provides an auditable, standardized baseline.
NIST AI RMF (AI Risk Management)	Four functions defined: Govern, Map, Measure, Manage.	HNS-36 coordinates + audit logs (spec §4.2) provide concrete measurement axes for the Measure function.

6.2 Roadmap to International Standardization

Stage 1: Pre-Standard Specification (Current)

GitHub specification documents and an eleven-volume ISBN series are publicly available. The machine-readable JSON-LD / RDF format requirement (spec §2.2) means the specification is already structured for programmatic verification. The first standardization prerequisite—public availability of a verifiable specification—is satisfied.

Stage 2: Independent Verification

Third-party verification of the HNS-36 matrix by research institutions and standards bodies. Cognitive-scientific assessment of the six-layer causal structure. Reproducibility confirmation through pilot implementations with JSON-LD validation tooling.

Stage 3: Pilot Certification

HNS-compliant pilot implementations in high-risk domains. Integration experiments with EU AI Act

conformity assessment. Development of reference JSON-LD / RDF schemas for public review.

Stage 4: Proposal to ISO/IEC JTC1/SC42

New Work Item Proposal positioning HNS-36 as a standard for cognitive coordinate systems for AI reasoning. The Tamper-Resistance Specification v1.0 becomes the basis for a hardware implementation certification scheme. The JSON-LD / RDF verification requirement becomes an annex to the main standard.

6.3 The Trust Infrastructure This Creates

If adopted as an international standard, the chain of consequence is:

- HNS-36 (JSON-LD / RDF encoded) becomes the externally verifiable criterion for AI safety evaluation.
- Third-party certification bodies verify compliance against the Tamper-Resistance Specification.
- Only certified systems may operate in high-risk domains: healthcare, finance, legal, critical infrastructure.
- HNS hardware with cryptographic attestation functions as a certificate of AI trustworthiness at the silicon level.
- The roles of SSL/TLS certificates, GMP certification, and aircraft type certification are brought to the AI industry.

Chapter 7 Industrial Applications

7.1 Edge AI: The Primary Deployment Context

Gartner projects that by 2027 SLM usage will surpass LLM usage, and 75% of enterprise data will be processed at the edge. The fundamental challenge for Edge AI is not model size but context retention and intent understanding. An HNS chip embedded in an edge device delivers four simultaneous properties grounded in the specification:

- Privacy: Confidential data stays on device via M5 local storage (HNS-864). Non-HNS semantic outputs are invalidated at the chip level (spec §3.3).
- Availability: Structural Integrity Monitoring and Non-Mapping Rule enforcement operate fully offline (spec §2.3, §2.1).
- Low latency: No cloud round-trip; coordinate assignment and consistency checks execute on the local HNS module.
- Uniformity: Because chip specification is physically fixed (ROM/ASIC), safety guarantees are identical across every device in a fleet—no configuration variance (spec §3.1).

7.2 Gateway AI: Routing for Task-Specific AI Ecosystems

As medical AI, legal AI, educational AI, manufacturing AI, and mental health AI proliferate, the challenge of selecting and orchestrating them becomes acute. The HNS Gateway AI architecture resolves this structurally: user input is converted to HNS coordinates, and the appropriate specialist model is determined automatically from the coordinate position. Each specialist receives only domain-internal inputs. The Non-Mapping Rule (spec §2.1) prevents cross-domain semantic contamination between specialists.

7.3 Priority Application Domains

Sector	HNS Hardware Value	Regulatory Alignment
Healthcare	Patient data retained on Edge via M5 (HNS-864). Structural Integrity Monitoring prevents diagnostic hallucination. Bypass Prevention ensures no data leaks through non-HNS paths.	HIPAA, Medical Device Regulations, Data Protection Laws
Financial Services	Transaction data never reaches the cloud. Category drift monitoring (spec §2.3) detects semantic anomalies in real-time fraud detection.	FSA Guidelines, Basel Accords, PCI DSS
Defense & Security	Air-gapped deployment with ROM/ASIC fixation. Version spoofing prohibition (spec §4.1) ensures coordinate system integrity cannot be faked.	Classified Info Protection Laws, NATO STANAG, Defense Procurement Standards
Manufacturing	Network-isolated AI with Bypass	IEC 62443, NIST SP 800-82 (OT

& OT	Prevention (spec §3.3) physically blocking external semantic modules from connecting to OT systems.	<i>Security)</i>
Care & Education	Individual profiles stored locally via HNS-864. Non-Mapping Rule prevents external service from reading or reconstructing personal semantic patterns.	<i>GDPR, Personal Data Protection Laws, Children's Online Privacy</i>

Chapter 8 Integrated Comparison: Cloud AI vs. HNS Hardware

The table below integrates the strategic analysis with specification requirements, providing a comprehensive comparison across nine dimensions.

Risk / Requirement	Cloud AI (Conventional)	HNS + Edge AI with Hardware (Spec v1.0)
Confidential data flow	Natural language transmitted to cloud servers	Only HNS coordinates transmitted; raw data stays on Edge (M5 local, HNS-864)
Resistance to tampering	Software updates allow modification at any time	ROM/FPGA/ASIC fixation; modification physically impossible (spec §3.1)
Proof of safety	Vendor self-assessment; no external verification	JSON-LD/RDF machine-verifiable structure + cryptographic attestation (spec §2.2, §3.3)
Hallucination prevention	No structural mechanism; output filtering only	Layer & Category drift monitoring + Non-Mapping Rule reject semantic drift at source (spec §2.3)
Prompt injection	Input passed directly to model; structurally exposed	External vocabulary injection treated as structural tampering; auto-rejected (spec §2.1, §2.4)
Bypass risk	No hardware isolation; software can route around guardrails	Physical bus isolation; non-HNS semantic output invalidated (spec §3.3)
Auditability	No structural log; black-box inference	Input–output coordinate mapping logged; tampering attempts recorded (spec §4.2)
Model dependency	Tied to specific vendor models	Model-agnostic; HNS chip operates independently of model processor

Chapter 9 Analysis and Conclusions

9.1 The Transformation: Assertion to Proof

The essential significance of implementing HNS in hardware is the transformation of AI safety from assertion to proof. Today's AI industry has vendors asserting safety but no structure for proving it. The integration of the Structural Integrity Specification with hardware implementation creates that structure through three converging principles:

- Verifiability: JSON-LD / RDF machine-readable structure enables programmatic third-party verification without access to model internals.
- Transparency: Continuous coordinate logging (Layer assignment, Category assignment, drift detection) creates an auditable structural record of every inference.

- Immutability: ROM / FPGA / ASIC fixation of the HNS-36 coordinate system, Non-Mapping Rule, and Bypass Prevention logic makes modification physically impossible.

This is equivalent to building a trust infrastructure for the AI industry—analogous to PKI for the internet, GMP for pharmaceuticals, and type certification for aviation.

9.2 Why Competing Organizations Cannot Arrive at This Point

HNS does not sell model AI. It provides a coordinate system and specification. The entity setting the standard is separate from the entities building products on that standard. This structural separation—which mirrors electrical safety marks (PSE, CE) and medical device ISO 13485 certification—is what makes the standard credible. No AI model vendor can achieve this through its own efforts: the conflict of interest is architectural, not ethical.

9.3 Remaining Technical Challenges

- HNS-36 coordinate transformation engine design: Low-power, low-latency ASIC implementation of coordinate assignment and consistency checking circuits.
- Model AI interface standardization: Communication protocol definition for the parallel-execution architecture between HNS chip and model processor.
- JSON-LD / RDF schema publication: Reference schemas enabling third-party tooling to validate HNS-36 cell definitions programmatically.
- Third-party cognitive-scientific validation: Independent assessment of the HNS-36 causal layer structure against established neuroscience and cognitive science findings.
- ISO/IEC JTC1/SC42 engagement: New Work Item Proposal preparation incorporating the Tamper-Resistance Specification as a hardware annex.
- Pilot certification: Proof-of-concept deployments in healthcare and legal services with full audit log generation.

9.4 Conclusion

The future of AI is at the edge. As Edge AI penetrates healthcare, elder care, education, manufacturing, and daily life, the decisive question shifts from "how capable is it?" to "how do we prove it can be trusted?"

HNS hardware implementation—grounded in the Structural Integrity and Tamper-Resistance Specification v1.0—provides the most structurally complete answer currently available. The Non-Mapping Rule blocks semantic invasion. The JSON-LD / RDF encoding enables machine verification. The ROM/ASIC fixation makes the coordinate system physically immutable. The audit log makes inference traceable. Together, these properties establish the external verification mechanism that genuine proof of AI safety has always required but never had.

From claiming to be safe, to structurally proving it. That is what HNS hardware implementation means.

References

Primary Sources

- Hara, S. HNS Structural Integrity and Tamper-Resistance Specification v1.0. Natural Structure Works / GitHub: satoru-hara/03_NSW, 2026.
- Natural Structure Works. Human Natural Structure Foundational Specification 1.0. NSW Official Website, 2026.
- Natural Structure Works. HumanOS Blueprint — Minimal Implementation Specification v1.0. GitHub: satoru-hara/03_NSW, 2026.
- Natural Structure Works. LCS-Large-Cognitive-Structure / The-Potential-of-HNS. GitHub: satoru-hara/03_NSW, 2026.
- Natural Structure Works. HNS-Series Preamble / HNS Structure / HNS Technical Note. GitHub: satoru-hara/03_NSW, 2026.

Standards and Regulatory Documents

- ISO/IEC 42001:2023. Information Technology — Artificial Intelligence — Management System. ISO, 2023.
- European Commission. EU AI Act (Regulation 2024/1689). Official Journal of the European Union, 2024.
- NIST. AI Risk Management Framework (AI RMF 1.0). National Institute of Standards and Technology, 2023.
- OWASP. OWASP Top 10 for Large Language Model Applications v2025. OWASP Foundation, 2025.
- IPA (Information-technology Promotion Agency, Japan). 10 Major Information Security Threats 2026, Organizational Edition. March 2026.

Industry Research and Analysis

- Gartner. Gartner Predicts SLM Usage Will Surpass LLM Usage by 2027. Gartner Research, 2025.
- Dell Technologies. The Power of Small: Edge AI Predictions for 2026. Dell Blog, January 2026.
- NVIDIA Research. Small Language Models are the Future of Agentic AI. NVIDIA Technical Report, June 2025.
- N-iX. Key Edge AI Trends Transforming Enterprise Tech in 2026. N-iX Blog, February 2026.
- Trend Micro. 2026 Corporate Security Threat Predictions Report. December 2025.

This report integrates: *HNS Hardware Analysis Report (Claude / Anthropic, April 2026)* + *HNS Structural Integrity and Tamper-Resistance Specification v1.0 (S. Hara / NSW, 2026)*

© 2026 Natural Structure Works. All rights reserved.